Movie Rating Prediction

1. Load the data from .dat files.

```
[ ] #JaiShreeRam
```

```
[3] import pandas as pd
    import numpy as np
    import matplotlib.pyplot as plt

    # Step 1: Load the data from the .dat files
    users_df = pd.read_csv('users.dat', sep='::', engine='python', header=None, names=['UserID', 'Gender', 'Age', 'Occupation', 'ZipCode'])
    ratings_df = pd.read_csv('ratings.dat', sep='::', engine='python', header=None, names=['UserID', 'MovieID', 'Rating', 'Timestamp'])
    movies_df = pd.read_csv('movies.dat', sep='::', engine='python', header=None, names=['MovieID', 'Title', 'Genres'],encoding='latin1')
```

```
[4] users_df.head()
```

	UserID	Gender	Age	Occupation	ZipCode
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

```
ratings_df.head()
```

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

```
[6] movies_df.head()
```

	MovieID	Title	Genres
0	1	Toy Story (1995)	AnimationIChildren'slComedy
1	2	Jumanji (1995)	AdventureIChildren'slFantasy
2	3	Grumpier Old Men (1995)	ComedyIRomance
3	4	Waiting to Exhale (1995)	ComedyIDrama
4	5	Father of the Bride Part II (1995)	Comedy

2. Data cleaning

```
[7] #Handling Missing Values

    missing_age_count = users_df['Age'].isnull().sum()
    if missing_age_count > 0:
        print(f"Number of missing values in 'Age': {missing_age_count}")
```

```
[8] users_df.drop_duplicates(subset=['UserID'], inplace=True)
    ratings_df.drop_duplicates(subset=['UserID', 'MovieID'], inplace=True)
```

3. Explore and visualize the data

```python
# Explore user age distribution
plt.hist(users_df['Age'], bins=20, color='skyblue', edgecolor='black')
plt.title('User Age Distribution')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```

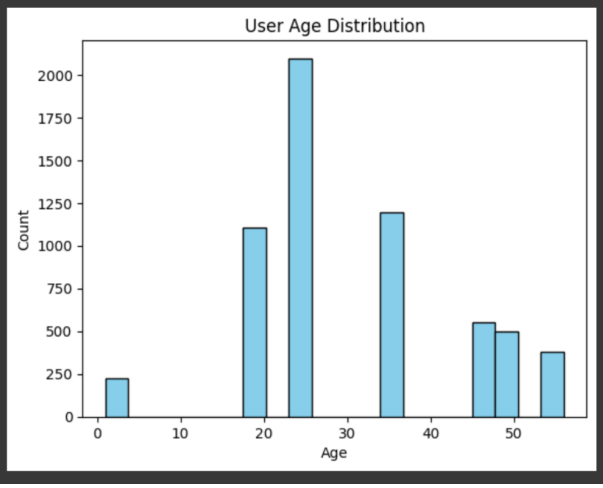


```python
# Explore movie ratings distribution
plt.hist(ratings_df['Rating'], bins=5, color='lightcoral', edgecolor='black', range=(0.5, 5.5))
plt.title('Movie Ratings Distribution')
plt.xlabel('Rating')
plt.ylabel('Count')
plt.xticks([1, 2, 3, 4, 5])
plt.show()
```

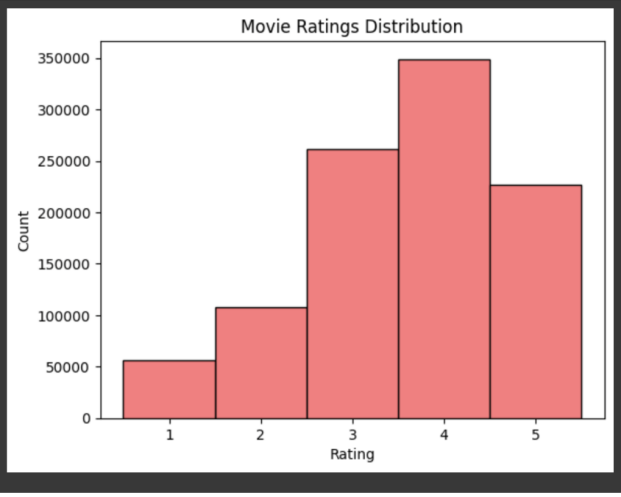

4. Split the dataset.

```python
[11] #Splitting the dataset

    from sklearn.model_selection import train_test_split

    test_size = 0.2

    ratings_train, ratings_test = train_test_split(ratings_df, test_size=test_size, random_state=42)
```

5. Model training and testing.

```
[12] !pip install scikit-surprise
```

```
Collecting scikit-surprise
  Downloading scikit-surprise-1.1.3.tar.gz (771 kB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 772.0/772.0 kB 7.5 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: joblib>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-surprise) (1.3.2)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from scikit-surprise) (1.23.5)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-surprise) (1.11.3)
Building wheels for collected packages: scikit-surprise
  Building wheel for scikit-surprise (setup.py) ... done
  Created wheel for scikit-surprise: filename=scikit_surprise-1.1.3-cp310-cp310-linux_x86_64.whl size=3163340 sha256=a4cf171b7046f74ce3372e
  Stored in directory: /root/.cache/pip/wheels/a5/ca/a8/4e28def53797fdc4363ca4af740db15a9c2f1595ebc51fb445
Successfully built scikit-surprise
Installing collected packages: scikit-surprise
Successfully installed scikit-surprise-1.1.3
```

```
from surprise import Dataset, Reader, SVD
from surprise.model_selection import train_test_split
from surprise import accuracy


reader = Reader(rating_scale=(1, 5))
data = Dataset.load_from_df(ratings_train[['UserID', 'MovieID', 'Rating']], reader)


model = SVD(reg_all=0.02)

trainset = data.build_full_trainset()
model.fit(trainset)


testset = ratings_test[['UserID', 'MovieID', 'Rating']].values.tolist()

predictions = model.test(testset)
```

```
[22] mae = accuracy.mae(predictions)
rmse = accuracy.rmse(predictions)

print(f"Mean Absolute Error (MAE): {mae:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")
```

```
MAE:  0.6885
RMSE: 0.8763
Mean Absolute Error (MAE): 0.6885
Root Mean Squared Error (RMSE): 0.8763
```

```
[23] user_id_to_predict = 6040
movie_id_to_predict = 1096


predicted_rating = model.predict(user_id_to_predict, movie_id_to_predict).est

print(f"Predicted Rating for User {user_id_to_predict} and Movie {movie_id_to_predict}: {predicted_rating:.2f}")
```

```
Predicted Rating for User 6040 and Movie 1096: 3.92
```

```
ratings_df
```

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291
...	...	...	...	...
1000204	6040	1091	1	956716541
1000205	6040	1094	5	956704887
1000206	6040	562	5	956704746
1000207	6040	1096	4	956715648
1000208	6040	1097	4	956715569

1000209 rows × 4 columns

The predicted rating is **3.9** and the actual rating is **4**.